

Lecture 08 Concurrency and Synchronization

Dr. Tushar, Mosaddek Hossain Kamal
Professor

Computer Science and Engineering, University of Dhaka,
BSc Third Year, Semester 2 (July – Dec), Academic Year: 2022

CSE3201: Operating Systems

November 21, 2024

Outline

- ① Learning Outcome
- ② Concurrency and Synchronization
- ③ Test-and-Set
 - Producer-Consumer Problem using Test-and-Set
- ④ Semaphores
 - Producer-Consumer Problem with Semaphore
- ⑤ Monitors
 - Monitor example
 - Producer-Consumer using monitors
- ⑥ Example: Classical Synchronization problems

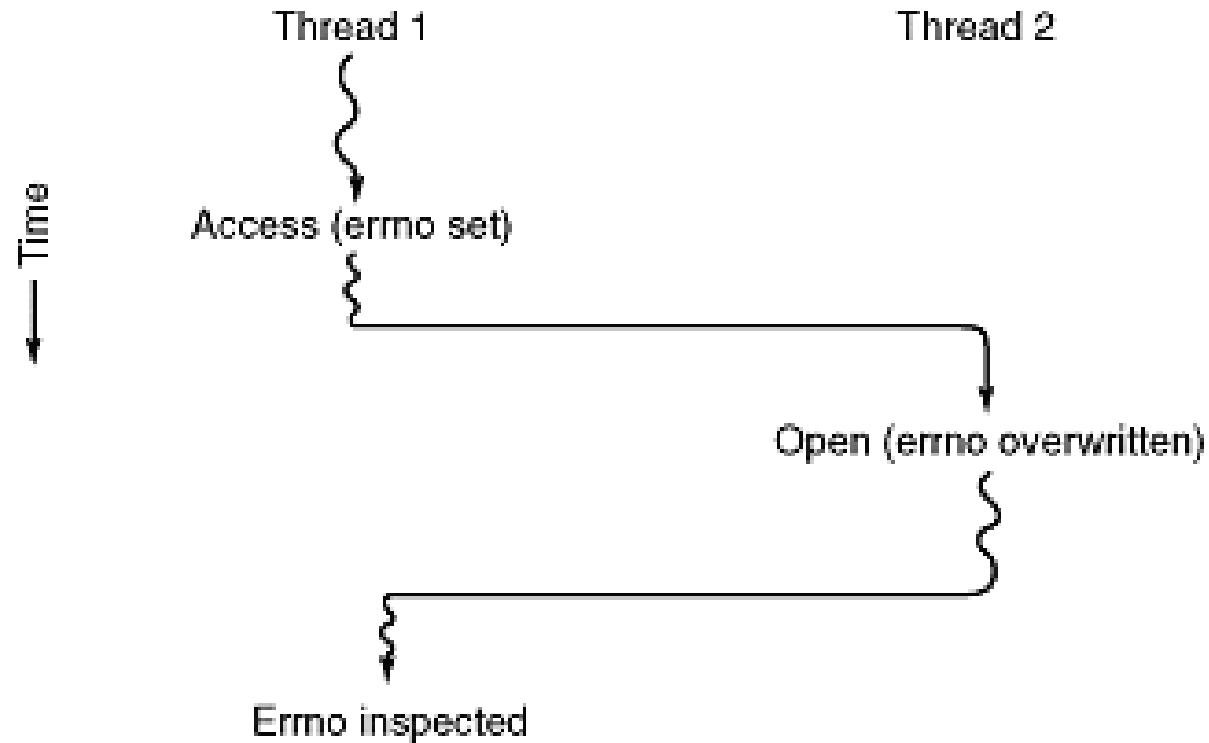
Learning Outcome

- **Understand the Operating System Process**
 - Process and Thread
 - Process and Thread context, data Structure
 - PCB and TCB
- **Understand the Process Management**
 - Context Switch and Dispatcher
 - Process Scheduling
 - Process and Thread Synchronization tools
 - Deadlock
- **Realtime process**
 - Concept of Soft and Hard real time process
 - Realtime Process Scheduling

Concurrency and Synchronization

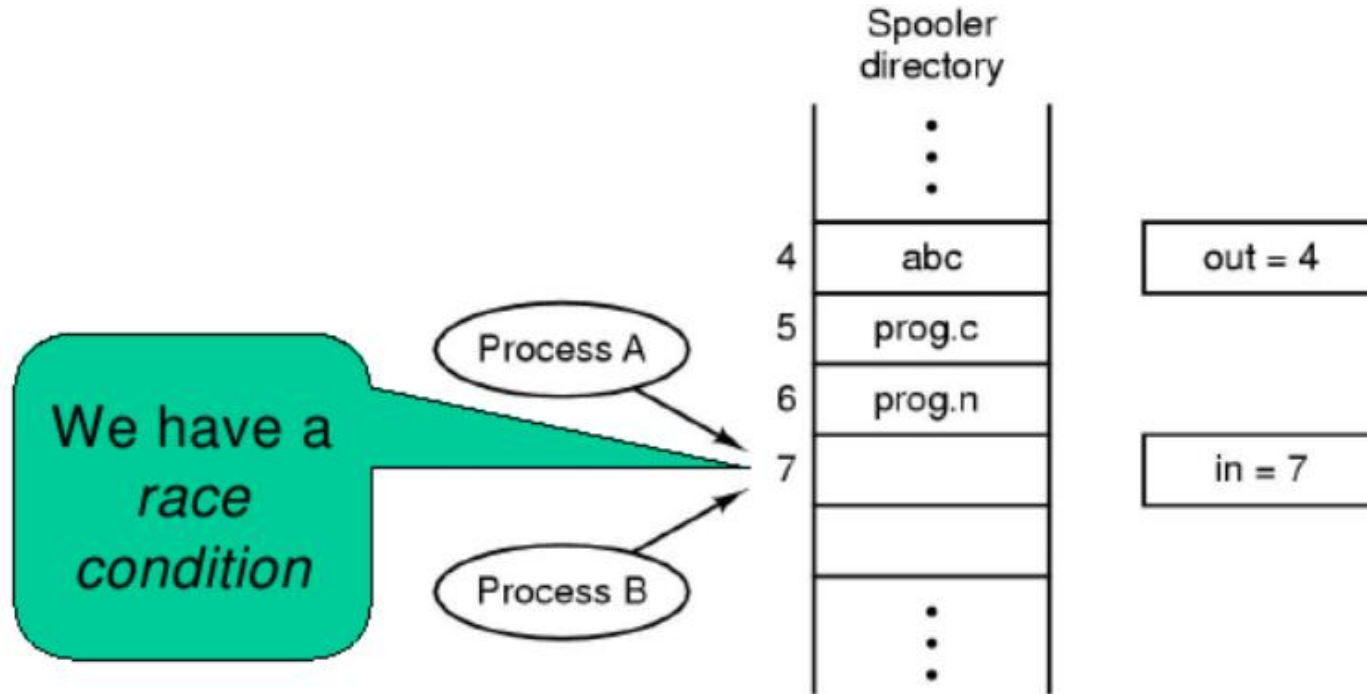
- Inter thread/Process communication
- Critical Region
- Solution
- Examples of IPC/ITC problems
- Tools used to overcome these IPC/ITC problems
 - Semaphore
 - Monitor
 - locks
- OS Synchronization Primitives (DUOS) – Cortex-M4
 - Lock
 - Semaphore
 - Condition variables

Making Single-Threaded Code Multithreaded



Conflict between threads over the use of a global variable

Inter-Thread and Process Communication



Two processes want to access shared memory at same time

Critical Region


- We can control access to the shared resources by controlling access to the code that accesses the resource

Critical Region

A critical region is a region of code where shared resources are accessed

- Variables, memory, files, etc.
- Unconditional entry to the critical region results in race condition
 - Incorrect behaviour, deadlock, lost work, ...

Critical Region



Mutual exclusion using critical regions

Example Critical Region

```
struct node {
    int data;
    struct node *next;
};
struct node *head;
void init(void)
{
    head = NULL;
}
```

Simple first-in-first-out
queue implemetation as a
linked list **Critical Region**

```
void insert(struct *item)
{
    item->next=head;
    head = item;
}
struct node *remove(void)
{
    struct node *t;
    t = head;
    if(t!=NULL){
        head = head->next;
    }
    return t;
}
```

Also known as critical section.

Conditions required of any solution to the critical region problem

- Mutual Exclusion:
 - No two processes simultaneously in critical region
- No assumptions made about speeds or numbers of CPUs
- Progress
 - No process running outside its critical region may block another process
- Bounded
 - No process must wait forever to enter its critical region

A solution?

- A lock variable
 - if lock == 1,
 - Somebody is in the critical section and we must wait
 - if lock == 0
 - nobody is in the critical section and we are free to enter

Solution

Process-0

```
while(TRUE) {  
    while(lock ==1);  
    lock=1;  
    critical();  
    lock=0;  
    non_critical();  
}
```

Process-1

```
while(TRUE) {  
    while(lock ==1);  
    lock=1;  
    critical();  
    lock=0;  
    non_critical();  
}
```

Problematic Execution Sequence

```
while(TRUE) {
```

```
    while(lock == 1);
```

```
    lock = 1;
```

```
    critical();
```

```
    lock = 0
```

```
    non_critical();
```

```
}
```

```
while(TRUE) {
```

```
    while(lock == 1);
```

```
    lock = 1;
```

```
    critical();
```

```
    lock = 0
```

```
    non_critical();
```

```
}
```



- Unfortunately, it is usually easier to show something does not work, than it is to prove that it does work.
 - Ideally, we'd like to prove, or at least informally demonstrate, that our solutions work.

Mutual Exclusion by Taking Turns

```
while (TRUE) {  
    while (turn != 0)    /* loop */ ;  
    critical_region();  
    turn = 1;  
    noncritical_region();  
}
```

(a)

```
while (TRUE) {  
    while (turn != 1)    /* loop */ ;  
    critical_region();  
    turn = 0;  
    noncritical_region();  
}
```

(b)

Proposed solution to critical region problem (a) Process-0, (b)
Process-1

Mutual Exclusion by Taking Turns

Works due to strict alternation

- Each process takes turns

Cons

- Busy waiting
- Process must wait its turn even while the other process is doing something else.
 - with many processes, must wait for everyone to have a turn
 - Does not guarantee progress if a process no longer needs a turn
 - Poor solution when processes require the critical section at differing rates

Mutual Exclusion by Disabling Interrupts

- Before entering a critical region, disable interrupts
- After leaving the critical region, enable interrupts
- Pros
 - simple
- Cons
 - Only available in the kernel
 - Blocks everybody else, even with no contention
 - Slows interrupt response time
 - Does not work on a multiprocessor

Hardware Support for mutual exclusion

Test and set instruction

- Can be used to implement lock variables correctly
 - it loads the value of the lock
 - if lock == 0
 - set the lock to 1
 - return result '0' – we acquire the lock
 - if lock == 1
 - return 1 – another thread/process has the lock
- Hardware guarantees that the instruction executes atomically
 - Atomically: As an indivisible unit.

Mutual Exclusion with Test-and-Set

enter_region:

```
TSL REGISTER,LOCK |copy lock to register and set lock to 1  
CMP REGISTER,#0    |was lock zero?  
JNE enter_region |if it was non zero, lock was set, so loop  
RET               |return to Caller; Critical region entered
```

leave_region:

```
MOVE LOCK,#0       |store a 0 in lock  
RET                |return to caller
```

Entering and leaving a critical region using the TSL instruction

Pros

- Simple (easy to show its correct)
- Available at user-level
 - To any number of processors
 - To implement any number of lock variables

Cons

- Busy waits (also termed a spin lock)
 - Consumes CPU
 - Livelock in the presence of priorities
 - If a low priority process has the lock and a high priority process attempts to get it, the high priority process will busy-wait forever.
 - Starvation is possible when a process leaves its critical section and more than one process is waiting.


Tackling the Busy-Wait Problem

Busy Waiting?? What can we do?

- Sleep / Wakeup
 - The idea when process is waiting for an event, it calls sleep to block. instead of busy waiting.
 - The the event happens, the event generator (another process) calls wakeup to unblock the sleeping process.

The Producer-Consumer Problem


- Also called the bounded buffer problem
- A producer produces data items and stores the items in a buffer
- A consumer takes the items out of the buffer and consumes them.



Issues

We must keep an accurate count of items in buffer

- Producer
 - can sleep when the buffer is full.
 - and wakeup when there is empty space in the buffer
 - The consumer can call wakeup when it consumes the first entry of the full buffer
- Consumer
 - Can sleep when the buffer is empty
 - And wake up when there are items available
 - Producer can call wakeup when it adds the first item to the buffer



Pseudo-code for producer and consumer

Producer

```
count = 0;
#define N 4 /* buf size */
prod(){
    while(TRUE) {
        item = produce()
        if(count == N)
            sleep();
        insert_item();
        count++;
        if(count == 1)
            wakeup(con);
    }
}
```

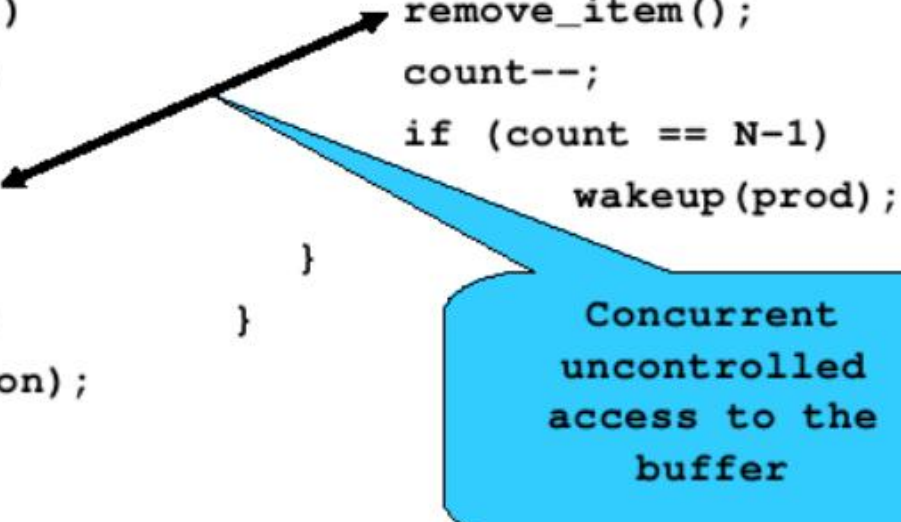
Consumer

```
con(){
    while(TRUE){
        if(count == 0)
            sleep();
        remove_item();
        count++;
        if(count == N-1)
            wakeup();
    }
}
```


Problems

```
int count = 0;
#define N 4 /* buf size */
prod() {
    while(TRUE) {
        item = produce()
        if (count == N)
            sleep();
        insert_item();
        count++;
        if (count == 1)
            wakeup(con);
    }
}

con() {
    while(TRUE) {
        if (count == 0)
            sleep();
        remove_item();
        count--;
        if (count == N-1)
            wakeup(prod);
    }
}
```

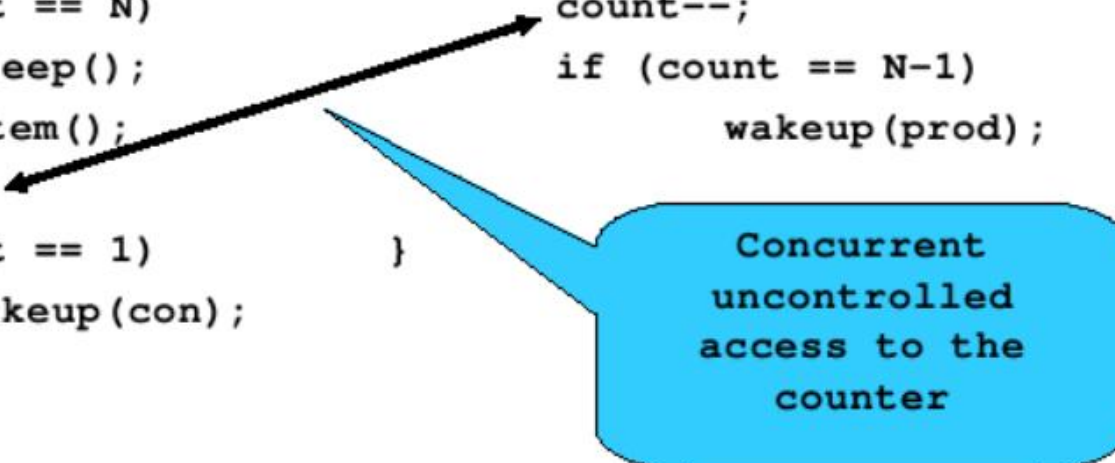


Concurrent uncontrolled access to the buffer

Problems

```
int count = 0;
#define N 4 /* buf size */
prod() {
    while(TRUE) {
        item = produce()
        if (count == N)
            sleep();
        insert_item();
        count++;
        if (count == 1)
            wakeup(con);
    }
}

con() {
    while(TRUE) {
        if (count == 0)
            sleep();
        remove_item();
        count--;
        if (count == N-1)
            wakeup(prod);
    }
}
```



Concurrent uncontrolled access to the counter

Proposed Solution


Lets use a locking primitive based on test-and-set to protect the concurrent access

```
int count = 0;
#define N 4 /* buf size */
prod() {
    while(TRUE) {
        item = produce()
        if (count == N)
            sleep();
        acquire_lock()
        insert_item();
        count++;
        release_lock()
        if (count == 1)
            wakeup(con);
    }
}

con() {
    while(TRUE) {
        if (count == 0)
            sleep();
        acquire_lock()
        remove_item();
        count--;
        release_lock();
        if (count == N-1)
            wakeup(prod);
    }
}
```

Problematic execution sequence

```
prod() {  
    while(TRUE) {  
        item = produce()  
        if (count == N)  
            sleep();  
        acquire_lock()  
        insert_item();  
        count++;  
        release_lock()  
        if (count == 1)  
            wakeup(con);  
    }  
}  
  
con() {  
    while(TRUE) {  
        if (count == 0)  
            sleep();  
        acquire_lock()  
        remove_item();  
        count--;  
        release_lock();  
        if (count == N-1)  
            wakeup(prod);  
    }  
}
```



wakeup without a matching sleep is lost

Problem

- The test for *some condition* and actually going to sleep needs to be atomic
- the following does not work

```
acquire_lock();  
if (count == N)  
    sleep();  
release_lock();
```

- The lock is held while asleep \Rightarrow count will never change

Semaphores

- **Dijkstra (1965) introduced two primitives that are more powerful than simple sleep and wakeup alone.**
 - P(): proberen, from Dutch to *test*.
 - V(): verhogen, from Dutch to *increment*.
 - Also called wait and signal , down & up

How it works?

- If a resource is not available, the corresponding semaphore blocks any process **waiting** for the resource
- Blocked processes are put into a process queue maintained by the semaphore (avoids busy waiting!)
- when a process releases a resource, it **signals** this by means of the semaphore
- Signalling resumes a blocked process if there is any
- Wait and signal operations cannot be interrupted
- Complex coordination can be implemented by multiple semaphores

Semaphore Implementation

- Define a semaphore as a record

```
typedef struct {  
    int count;  
    struct process *L; //PCB/TCB listed  
}semaphore;
```

- Assume two simple operations:
 - **sleep** suspends the process that invokes it.
 - **wakeup(P)** resumes the execution of a blocked process P.

Semaphore operations now defined as

```
wait(S):  
    S.count--;  
    if(S.count<0){  
        add this process to S.L;  
        sleep  
    }  
  
signal(S):  
    S.count++;  
    if(S.count<=0){  
        remove a process P from S.L;  
        wakeup(P);  
    }
```

- Each primitive is atomic

Semaphore as a general sechronization tools

- Execute B in P_j only after A Executed in P_i
- Use semaphore *count* initialized to 0
- code:

P_i	P_j
\vdots	\vdots
A	$wait(flag)$
$signal(flag)$	B

Semaphore implementation of a Mutex

- Mutex also called for Mutual Exclusion
 - Can also be called a lock

```
semaphore mutex;
```

```
mutex.count = 1; /*initialize mutex */
```

```
wait(mutex); /*before enter the critical region */
```

```
BlahBlah ..
```

```
signal(mutex); /* exit the critical region */
```

Notice that the initial count determines how many waits can progress before blocking and requiring a signal \Rightarrow mutex.count initialized as 1

Solving the producer-consumer problem with semaphores

```
#define N = 4

semaphore mutex = 1;

/* count empty slots */
semaphore empty = N;

/* count full slots */
semaphore full = 0;
```

Solving the producer-consumer problem with semaphore

```
prod() {  
    while(TRUE) {  
        item = procedure();  
        wait(empty);  
        wait(mutex);  
        insert_item();  
        signal (mutex);  
        signal(full);  
    }  
}  
  
con() {  
    while(TRUE) {  
        wait(full);  
        wait(mutex);  
        remove_item;  
        signal (mutex);  
        signal(empty);  
    }  
}
```

Summarizing Semaphore


- Semaphore can be used to solve a variety of concurrency problems
- However, programming with semaphore can be error-prone
 - E.g. must signal for every wait for mutexes
 - Too many, or too few signals or waits, or signals and waits in the wrong order, can have catastrophic results

Monitors

- To ease the concurrency programming, Hoare (1974) proposed *monitors*
 - A higher level synchronization primitive
 - Programming language construct
- Idea
 - A set of procedure, variables, data types are grouped in a special kind of module, a monitor.
 - Variable and data types are only accessed from within the monitor
 - Only one process/thread can be in the monitor at any one time
 - Mutual execution is implemented by the compiler (which should be less error-prone)

Monitor

- When a thread calls a monitor procedure that has a thread already inside, it is queued and it sleeps until the current thread exits the monitor



Monitors

Example of a monitor:

```
monitor example
    integer i;
    condition c;

    procedure producer();
    .
    .
    .
    end;

    procedure consumer();
    .
    .
    .
    end;
end monitor
```


Simple Example

```
monitor counter {  
    int count;  
    procedure inc(){  
        count = count + 1;  
    }  
    procedure dec(){  
        count = count - 1;  
    }  
}
```

Note: “paper” language

- Compile guarantees only one thread can be active in the monitor at any one time
- Easy to see this provides mutual exclusion
 - No race condition


How do we block waiting for a event?

- We need a mechanism to block waiting for an event (in addition to ensure mutual exclusion)
 - e.g., for producer consumer problem when buffer is empty or full
- *Conditon Variables*

Condition Variable

- To allow a process to wait within the monitor, a **condition** variable must be declared as
condition x, y
- Condition variable can only be used with the operations **wait** and **signal**
 - The operation **x.wait();**
means that the process invoking this operation is suspended until another process invokes **x.signal();**
 - The **x.signal** operation resumes exactly one suspended process. If no process is suspended, then the **signal** operation has no effect.

Condition Variables



Monitors

```
monitor ProducerConsumer
  condition full, empty;
  integer count;
  procedure insert(item: integer);
  begin
    if count = N then wait(full);
    insert_item(item);
    count := count + 1;
    if count = 1 then signal(empty)
  end;
  function remove: integer;
  begin
    if count = 0 then wait(empty);
    remove = remove_item;
    count := count - 1;
    if count = N - 1 then signal(full)
  end;
  count := 0;
end monitor;
```

```
procedure producer;
begin
  while true do
    begin
      item = produce_item;
      ProducerConsumer.insert(item)
    end
  end;
procedure consumer;
begin
  while true do
    begin
      item = ProducerConsumer.remove;
      consume_item(item)
    end
  end;
end;
```


Outline of producer-consumer problem with monitors

- only one monitor procedure active at one time
- buffer has N slots

Example: Some Classical Problems

Dining Philosophers Problem

- Philosopher eat/think
- Eating needs 2 forks
- Pick one fork at a time
- How to prevent deadlock



One Dining-Philosopher problem Solution

```
#define N          5          /* number of philosophers */
#define LEFT      (i+N-1)%N   /* number of i's left neighbor */
#define RIGHT     (i+1)%N     /* number of i's right neighbor */
#define THINKING  0          /* philosopher is thinking */
#define HUNGRY    1          /* philosopher is trying to get forks */
#define EATING    2          /* philosopher is eating */
typedef int semaphore;        /* semaphores are a special kind of int */
int state[N];                /* array to keep track of everyone's state */
semaphore mutex = 1;         /* mutual exclusion for critical regions */
semaphore s[N];              /* one semaphore per philosopher */

void philosopher(int i)      /* i: philosopher number, from 0 to N-1 */
{
    while (TRUE) {           /* repeat forever */
        think( );            /* philosopher is thinking */
        take_forks(i);        /* acquire two forks or block */
        eat( );               /* yum-yum, spaghetti */
        put_forks(i);         /* put both forks back on table */
    }
}
```

Solution to dining philosophers problem (part 1)

One Dining-Philosopher problem Solution contd.

```
#define N 5                                /* number of philosophers */

void philosopher(int i)                    /* i: philosopher number, from 0 to 4 */
{
    while (TRUE) {
        think( );                          /* philosopher is thinking */
        take_fork(i);                      /* take left fork */
        take_fork((i+1) % N);              /* take right fork; % is modulo operator */
        eat( );                            /* yum-yum, spaghetti */
        put_fork(i);                       /* put left fork back on the table */
        put_fork((i+1) % N);               /* put right fork back on the table */
    }
}
```

A nonsolution to dining philosophers problem

Dining Philosopher Problem Solution

```
void take_forks(int i)                /* i: philosopher number, from 0 to N-1 */
{
    down(&mutex);                     /* enter critical region */
    state[i] = HUNGRY;                /* record fact that philosopher i is hungry */
    test(i);                          /* try to acquire 2 forks */
    up(&mutex);                       /* exit critical region */
    down(&s[i]);                       /* block if forks were not acquired */
}


void put_forks(i)                     /* i: philosopher number, from 0 to N-1 */
{
    down(&mutex);                     /* enter critical region */
    state[i] = THINKING;              /* philosopher has finished eating */
    test(LEFT);                       /* see if left neighbor can now eat */
    test(RIGHT);                      /* see if right neighbor can now eat */
    up(&mutex);                       /* exit critical region */
}

void test(i)                          /* i: philosopher number, from 0 to N-1 */
{
    if (state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {
        state[i] = EATING;
        up(&s[i]);
    }
}
```

Solution to dining philosophers problem (part 2)

Readers and Writers Problem

- Models access to a database
 - E.g., airline reservation system
- Can have more than one concurrent reader
 - To check schedules and reservations
- Writers must have exclusive access
 - To book a ticket or update a schedule



Readers and Writers Problem

```
typedef int semaphore;          /* use your imagination */
semaphore mutex = 1;           /* controls access to 'rc' */
semaphore db = 1;              /* controls access to the database */
int rc = 0;                    /* # of processes reading or wanting to */

void reader(void)
{
    while (TRUE) {              /* repeat forever */
        down(&mutex);           /* get exclusive access to 'rc' */
        rc = rc + 1;            /* one reader more now */
        if (rc == 1) down(&db); /* if this is the first reader ... */
        up(&mutex);             /* release exclusive access to 'rc' */
        read_data_base();       /* access the data */
        down(&mutex);           /* get exclusive access to 'rc' */
        rc = rc - 1;            /* one reader fewer now */
        if (rc == 0) up(&db);    /* if this is the last reader ... */
        up(&mutex);             /* release exclusive access to 'rc' */
        use_data_read();        /* noncritical region */
    }
}

void writer(void)
{
    while (TRUE) {              /* repeat forever */
        think_up_data();        /* noncritical region */
        down(&db);              /* get exclusive access */
        write_data_base();      /* update the data */
        up(&db);                /* release exclusive access */
    }
}
```

A solution to the readers and writers problem

Sleeping Barber Problem



Figure 1: Sleeping Barber Problem

Sleeping Barber Problem

```
#define CHAIRS 5                /* # chairs for waiting customers */

typedef int semaphore;          /* use your imagination */

semaphore customers = 0;        /* # of customers waiting for service */
semaphore barbers = 0;         /* # of barbers waiting for customers */
semaphore mutex = 1;           /* for mutual exclusion */
int waiting = 0;               /* customers are waiting (not being cut) */

void barber(void)
{
    while (TRUE) {
        down(&customers);        /* go to sleep if # of customers is 0 */
        down(&mutex);            /* acquire access to 'waiting' */
        waiting = waiting - 1;    /* decrement count of waiting customers */
        up(&barbers);            /* one barber is now ready to cut hair */
        up(&mutex);              /* release 'waiting' */
        cut_hair();              /* cut hair (outside critical region) */
    }
}

void customer(void)
{
    down(&mutex);                /* enter critical region */
    if (waiting < CHAIRS) {      /* if there are no free chairs, leave */
        waiting = waiting + 1;  /* increment count of waiting customers */
        up(&customers);         /* wake up barber if necessary */
        up(&mutex);             /* release access to 'waiting' */
        down(&barbers);         /* go to sleep if # of free barbers is 0 */
        get_haircut();          /* be seated and be serviced */
    } else {
        up(&mutex);             /* shop is full; do not wait */
    }
}
```

See the textbook